

CS 4F03 – Distributed Systems:

Final Project Report

Stuart Douglas – 1214422
Matthew Pagnan – 1208693

April 10, 2016

Contents

1	Description of Parallelization	3
1.1	Overview	3
1.2	Vector3D	3
1.3	OpenACC Data	3
2	Computing Parameters for Frames	3
3	Mandelbox Performance vs. Mandelbulb	4
4	Source Code	6
4.1	getcolor.cc	7
4.2	main.cc	8
4.3	raymarching.cc	11
4.4	renderer.cc	14
4.5	vector3d.h	16
4.6	3d.cc	18
4.7	3d.h	25
5	Running the Program	26
6	Bonus Features	27
6.1	Automatic Navigation	27

Video

The final video is 4 minutes at 30 FPS, each frame being 1920 x 1080 resolution. It can be found at <https://youtu.be/bwmpGCiiVN4>.

1 Description of Parallelization

1.1 Overview

The parallelization of the project is based off of the nested for-loop in `renderFractal` that iterates over each pixel in the output image, calculating the correct colour for that pixel based off of the passed parameters. When the loop is encountered, the necessary data from the CPU's memory is copied in. A kernel is then launched on the GPU, and each thread begins the processing for one pixel. Every function call within the kernel is run sequentially on the calling thread, with all subsequent routines from *those* functions inlined. Finally, the image data is copied back to the CPU.

1.2 Vector3D

Originally, a `vec3` was represented by a C++ class. We changed this implementation to a struct, and wrote macros to perform computations on the vector. There were a few macros originally included with the project, but many more had to be written to ensure all `vec3` operations could be run on the `vec3` struct. Using macros simplified integration with OpenACC, as routines are a relatively new feature and not fully robust yet.

1.3 OpenACC Data

The complex data that is private to each thread, such as the pixel data objects, the vectors storing the `colour` and the `to` vector for the pixel, and the double array containing the `farPoint` for the pixel are all stored in arrays, where each thread on the GPU accesses one element of each array. These values are not needed by the CPU at all, so the arrays are allocated on GPU memory using `acc_malloc`, and declared as device pointers. Note that they are allocated once at the beginning of the program and freed just before the program exits.

2 Computing Parameters for Frames

Generating the camera parameters for the next frame is done automatically, based on the furthest point in the mandelbox. To do this, additional data needs to be stored about each point. Each GPU thread will know how far the point is that it hits from `rayMarch`, as well as the vector representing that point. After the kernel has finished executing, we need to find the furthest of such points. Initially, it seems intuitive to use an OpenACC `max` reduction to find this distance as all the distance data is stored on the GPU, but we need to know the vector associated with that maximum distance, which is not supported by a reduction. Instead, we copy an array storing the distances to each point back from the GPU when the kernel exits, and the CPU iterates through them to find the max. Once it does, we use

`acc_memcpy_from_device` to copy the single vector associated with that distance back, without copying the entire array of pixel data. This does reduce performance a little, but the difference is relatively minor.

After the new vector is found, execution returns to `main`, and the CPU calculates the camera position and target for the next frame. The camera always moves directly towards the currently set “furthest point”, which may in fact be different than the camera’s target. This is because the camera’s target does not instantly snap to the new furthest point as soon as it’s set. Instead the camera interpolates between where it was looking before and where it “should” be looking. Even though it’s not looking directly at the new target, we still want it to follow the path towards that target, so as to prevent collisions. This essentially allows the camera to pan to face the new target so viewers can see not only the destination but they also get to see the environment around the camera. The camera will move towards the target until one of two things happen:

- The camera gets within a certain distance of the target point
- The camera is looking directly at the target

The first condition is very obvious for why the camera will change targets, as we don’t want to run into a wall. The second condition provides a more interesting video, by causing more camera movements rather than just locking on to a point and travelling in a straight line. That being said, it is undesirable to have the camera changing targets all the time, as the video would be extremely shaky. This is addressed by only allowing targets to change in the above two scenarios.

3 Mandelbox Performance vs. Mandelbulb

The reason the mandelbulb is faster to compute comes down to the distance estimator for mandelbulbs. Although we did not run the mandelbulb code with OpenACC (gpu servers were all in use), we can observe why it will be faster. Following is the distance estimator code for the mandelbox:

```
inline double MandelBoxDE(const vec3 &p0, const MandelBoxParams &params,
    double c1, double c2)
{
    vec3 p = p0;
    double rMin2 = SQR(params.rMin);
    double rFixed2 = SQR(params.rFixed);
    double escape = SQR(params.escape_time);
    double dfactor = 1;
    double r2 = -1;
    const double rFixed2rMin2 = rFixed2/rMin2;

    int i = 0;
    while (i < params.num_iter && r2 < escape)
    {
        COMPONENT_FOLD(p.x);
        COMPONENT_FOLD(p.y);
        COMPONENT_FOLD(p.z);
```

```

DOT_ASSIGN(r2,p);

if (r2 < rMin2)
{
    MULT_SCALAR(p, p, rFixed2rMin2);
    dfactor *= rFixed2rMin2;
}
else if (r2 < rFixed2)
{
    const double t = (rFixed2/r2);
    MULT_SCALAR(p, p, (rFixed2/r2));
    dfactor *= t;
}

dfactor = dfactor * fabs(params.scale) + 1.0;
MULT_SCALAR(p, p, params.scale);
ADD_POINT(p, p, p0);
i++;
}

return (MAGNITUDE(p) - c1) / dfactor - c2;
}

```

And for the mandelbulb:

```

inline double MandelBulbDE(const vec3 &p0, const MandelBoxParams &params)
{
    vec3 z;
    z = p0;

    double dr = 1.0;
    double r = 0.0;

    double Bailout = params.rMin;
    double Power = params.rFixed;

    for (int i = 0; i < params.num_iter; i++)
    {
        r = MAGNITUDE(z);
        if (r > Bailout) {
            break;
        }

        double theta = acos(z.z/r);
        double phi = atan2(z.y, z.x);
        dr = pow(r, Power - 1.0) * Power * dr + 1.0;

        double zr = pow(r, Power);
        theta = theta * Power;
        phi = phi * Power;

        z.x = zr*sin(theta)*cos(phi);
        z.y = zr*sin(phi)*sin(theta);
        z.z = zr*cos(theta);

        z.x = z.x + p0.x;
    }
}

```

```

        z.y = z.y + p0.y;
        z.z = z.z + p0.z;
    }
    return 0.5*log(r)*r/dr;
}

```

The mandelbox distance estimator does require two parameters that the mandelbulb does not, each of which need to be computed (`c1` and `c2`). But the main reason for the performance difference is that the mandelbulb distance estimator loop in general breaks out earlier than the box when `r > Bailout`. This is due to the nature of the computations in each distance estimator, and how they affect the value of `r` and `r2` for the mandelbulb and mandelbox, respectively. Additionally, the mandelbox distance estimator breaks out when the dot product of `p` is less than the square of the escape time. This compares to the mandelbulb, which checks the magnitude of `z` vs the minimum `r` value from the parameters file.

4 Source Code

The following source code files were not changed from the serial version of the program, and as such will their contents will not be reproduced in this document.

- `camera.h`
- `color.h`
- `getparams.c`
- `init3D.cc`
- `mandelbox.h`
- `renderer.h`
- `savebmp.c`

Due to OpenACC requirements for nested inline routines being in the same source file, the `MandelBoxDE` and `DE` functions were moved to `raymarching.cc`. Functions for printing progress and timing data were removed, as they were no longer deemed necessary due to the speedups from running the program on the GPU. As such, the following source code files were removed.

- `distance_est.cc`
- `mandelboxde.cc`
- `print.c`
- `timing.c`

Following are the modified source code files. Note that we made every effort to write “self-documenting” code by using clear variable names, so tried to limit comments within the source code to situations that are not immediately clear. Files are sorted alphabetically.

4.1 getcolor.cc

```
/*
This file is part of the Mandelbox program developed for the course
CS/SE Distributed Computer Systems taught by N. Nedialkov in the
Winter of 2015-2016 at McMaster University.

Copyright (C) 2015-2016 T. Gwosdz and N. Nedialkov

This program is free software: you can redistribute it and/or modify
it under the terms of the GNU General Public License as published by
the Free Software Foundation, either version 3 of the License, or
(at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
GNU General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program. If not, see <http://www.gnu.org/licenses/>.
*/

#include "color.h"
#include "renderer.h"
#include "vector3d.h"

//---lightning and colouring-----
#define CAM_LIGHT_W 1.8
#define CAM_LIGHT_MIN 0.3

#define CAM_LIGHT 1.0
#define BASE_COLOR 1.0
#define BACK_COLOR 0.4

inline void lighting(const vec3 &n, const vec3 &color, const vec3 &pos, const
    vec3 &direction, vec3 &outV)
{
    vec3 nn;
    SUBTRACT_SCALAR(nn, n, 1.0);
    double d = DOT(direction, nn);
    double ambient = MAX(CAM_LIGHT_MIN, d) * CAM_LIGHT_W;
    vec3 camLight;
    VEC(camLight, CAM_LIGHT, CAM_LIGHT, CAM_LIGHT);
    MULT_SCALAR(nn, camLight, ambient);
    MULT_POINTWISE(outV, color, nn);
}

#pragma acc routine seq
inline void getColour(vec3 &hitColor, const pixelData &pixData, const
    RenderParams &render_params, const vec3 &from, const vec3 &direction)
{
    VEC(hitColor, BASE_COLOR, BASE_COLOR, BASE_COLOR);

    if (pixData.escaped == false)
    {

```

```

//apply lighting
lighting(pixData.normal, hitColor, pixData.hit, direction, hitColor);

//add normal based colouring
if (render_params.colourType == 0 || render_params.colourType == 1) {
    MULT_POINTWISE(hitColor, hitColor, pixData.normal);
    ADD_SCALAR(hitColor, hitColor, 1.0);
    DIV_SCALAR(hitColor, hitColor, 2.0);
    MULT_SCALAR(hitColor, hitColor, render_params.brightness);

    //gamma correction
    clamp(hitColor, 0.0, 1.0);
    MULT_POINTWISE(hitColor, hitColor, hitColor)
}
if (render_params.colourType == 1)
{
    // "swap" colors
    double t = hitColor.x;
    hitColor.x = hitColor.z;
    hitColor.z = t;
}
}
else
{
    //we have the background colour
    VEC(hitColor, BACK_COLOR, BACK_COLOR, BACK_COLOR);
}
}

```

4.2 main.cc

```

/*
This file is part of the Mandelbox program developed for the course
CS/SE Distributed Computer Systems taught by N. Nedialkov in the
Winter of 2015-2016 at McMaster University.

Copyright (C) 2015-2016 T. Gwosdz and N. Nedialkov

This program is free software: you can redistribute it and/or modify
it under the terms of the GNU General Public License as published by
the Free Software Foundation, either version 3 of the License, or
(at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
GNU General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program. If not, see <http://www.gnu.org/licenses/>.
*/
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <sys/types.h>
#include "openacc.h"

```



```

#include "camera.h"
#include "renderer.h"
#include "mandelbox.h"
#include "vector3d.h"
#include "color.h"

#define NUM_FRAMES 7200
#define MOV_SPEED 0.001
#define CAM_SPEED 0.04
#define REACH_TARGET_THRESHOLD 0.5
#define NEW_TARGET_THRESHOLD 5

void getParameters(char *filename, CameraParams *camera_params, RenderParams
    *renderer_params,
                    MandelBoxParams *mandelBox_paramsP);
void init3D        (CameraParams *camera_params, const RenderParams *
    renderer_params);
void renderFractal(const CameraParams camera_params, const RenderParams
    renderer_params, unsigned char* image);
void saveBMP       (const char* filename, const unsigned char* image, int
    width, int height);

void genNewCamParams(CameraParams &curCam, CameraParams &nextCam){
    curCam.camPos[0] += (nextCam.camPos[0] - curCam.camPos[0])*0.01;
    curCam.camPos[1] += (nextCam.camPos[1] - curCam.camPos[1])*0.01;
    curCam.camPos[2] += (nextCam.camPos[2] - curCam.camPos[2])*0.01;
}

/*****
* Global Variables (public)
*****/
MandelBoxParams mandelBox_params;

// set by renderer.cc after each renderFractal
vec3 newFurthestPoint;

// device pointers, we allocate once at program start and free just before
// program exits
vec3* d_to;
vec3* d_colours;
double* d_farPoints;
pixelData* d_pixData;
double* d_distances;
#pragma acc declare deviceptr(d_to, d_colours, d_farPoints, d_pixData,
    d_distances) copyin(mandelBox_params)

int main(int argc, char** argv)
{
    CameraParams camera_params;
    RenderParams renderer_params;
    getParameters(argv[1], &camera_params, &renderer_params, &mandelBox_params)
        ;

    int image_size = renderer_params.width * renderer_params.height;

```

```

unsigned char *image = (unsigned char*)malloc(3*image_size*sizeof(unsigned
char));

d_to = (vec3*)acc_malloc(image_size * sizeof(vec3));
d_colours = (vec3*)acc_malloc(image_size * sizeof(vec3));
d_farPoints = (double*)acc_malloc(image_size * 3 * sizeof(double));
d_pixData = (pixelData*)acc_malloc(image_size * sizeof(pixelData));
d_distances = (double*)acc_malloc(image_size * sizeof(double));

vec3 furthestPoint;
char new_file_name[80];

// store generated images in their own directory
mkdir("images", 0777);

for (int i = 0; i < NUM_FRAMES; i++) {
    init3D(&camera_params, &renderer_params);
    renderFractal(camera_params, renderer_params, image);

    vec3 camTarget, camPos;
    VEC(camTarget, camera_params.camTarget[0], camera_params.camTarget[1],
camera_params.camTarget[2]);
    VEC(camPos, camera_params.camPos[0], camera_params.camPos[1],
camera_params.camPos[2]);

    if (i % 50 == 0) {
        printf("Done rendering frame %d. camTarget=[%f,%f,%f], camPos = [%f,%f
,%f], furthestPoint=[%f,%f,%f], newFurthestPoint=[%f,%f,%f]\n",
i, camTarget.x, camTarget.y, camTarget.z, camPos.x, camPos.y, camPos.z
, furthestPoint.x, furthestPoint.y, furthestPoint.z,
newFurthestPoint.x, newFurthestPoint.y, newFurthestPoint.z);
    }

    // only change target when:
    // - we have arrived at the point we were tracking
    // - we've finished locking on to the current target
    if (DISTANCE_APART(camPos, furthestPoint) < REACH_TARGET_THRESHOLD
|| DISTANCE_APART(camTarget, furthestPoint) < REACH_TARGET_THRESHOLD) {
        furthestPoint = newFurthestPoint;
    }

    // camera gradually points to the target
    camera_params.camTarget[0] += (furthestPoint.x - camTarget.x)*CAM_SPEED;
    camera_params.camTarget[1] += (furthestPoint.y - camTarget.y)*CAM_SPEED;
    camera_params.camTarget[2] += (furthestPoint.z - camTarget.z)*CAM_SPEED;

    // camera moves towards the target
    camera_params.camPos[0] += (furthestPoint.x - camPos.x)*MOV_SPEED;
    camera_params.camPos[1] += (furthestPoint.y - camPos.y)*MOV_SPEED;
    camera_params.camPos[2] += (furthestPoint.z - camPos.z)*MOV_SPEED;

    sprintf(new_file_name, "images/image_%d.bmp", i);
    saveBMP(new_file_name, image, renderer_params.width, renderer_params.
height);
}

```

```

    free(image);

    acc_free(d_to);
    acc_free(d_colours);
    acc_free(d_farPoints);
    acc_free(d_pixData);
    acc_free(d_distances);

    return 0;
}

```

4.3 raymarching.cc

```

/*
This file is part of the Mandelbox program developed for the course
CS/SE Distributed Computer Systems taught by N. Nedialkov in the
Winter of 2015-2016 at McMaster University.

Copyright (C) 2015-2016 T. Gwosdz and N. Nedialkov

This program is free software: you can redistribute it and/or modify
it under the terms of the GNU General Public License as published by
the Free Software Foundation, either version 3 of the License, or
(at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
GNU General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program. If not, see <http://www.gnu.org/licenses/>.
*/

#include "color.h"
#include "renderer.h"
#include "mandelbox.h"
#include "vector3d.h"

#define SQR(x) ((x)*(x))
#define COMPONENT_FOLD(x) { (x) = (fabs(x) <= 1) ? (x) : (((x) > 0) ? (2-(x)) : (-2-(x))); }

// #pragma acc declare copyin(mandelBox_params)
extern MandelBoxParams mandelBox_params;

extern double* d_distances;

inline double MandelBoxDE(const vec3 &p0, const MandelBoxParams &params,
    double c1, double c2)
{
    vec3 p = p0;
    double rMin2 = SQR(params.rMin);
    double rFixed2 = SQR(params.rFixed);
    double escape = SQR(params.escape_time);
    double dfactor = 1;

```

```

double r2          = -1;
const double rFixed2rMin2 = rFixed2/rMin2;

int i = 0;
while (i < params.num_iter && r2 < escape)
{
    COMPONENT_FOLD(p.x);
    COMPONENT_FOLD(p.y);
    COMPONENT_FOLD(p.z);

    DOT_ASSIGN(r2,p);

    if (r2 < rMin2)
    {
        MULT_SCALAR(p, p, rFixed2rMin2);
        dfactor *= rFixed2rMin2;
    }
    else if (r2 < rFixed2)
    {
        const double t = (rFixed2/r2);
        MULT_SCALAR(p, p, (rFixed2/r2));
        dfactor *= t;
    }

    dfactor = dfactor * fabs(params.scale) + 1.0;
    MULT_SCALAR(p, p, params.scale);
    ADD_POINT(p, p, p0);
    i++;
}

return (MAGNITUDE(p) - c1) / dfactor - c2;
}

inline double DE(const vec3 &p)
{
    double c1 = fabs(mandelBox_params.scale - 1.0);
    double c2 = pow(fabs(mandelBox_params.scale), 1 - mandelBox_params.
        num_iter);
    double d = MandelBoxDE(p, mandelBox_params, c1, c2);

    return d;
}

inline void normal(const vec3 & p, vec3 & normal)
{
    // compute the normal at p
    const double sqrt_mach_eps = 1.4901e-08;
    double eps = MAX( MAGNITUDE(p), 1.0 ) * sqrt_mach_eps;
    vec3 t1, e1;
    double x;

    VEC(e1, eps, 0, 0);
    ADD_POINT(t1, p, e1);
    x = DE(t1);
    SUBTRACT_POINT(t1, p, e1);
    normal.x = x - DE(t1);
}

```

```

    VEC(e1, 0, eps, 0);
    ADD_POINT(t1, p, e1);
    x = DE(t1);
    SUBTRACT_POINT(t1, p, e1);
    normal.y = x - DE(t1);

    VEC(e1, 0, 0, eps);
    ADD_POINT(t1, p, e1);
    x = DE(t1);
    SUBTRACT_POINT(t1, p, e1);
    normal.z = x - DE(t1);

    NORMALIZE(normal);
}

#pragma acc declare copyin(mandelBox_params)
#pragma acc routine seq
void rayMarch(const RenderParams &render_params, const vec3 &from, const vec3
    &direction, double eps,
    pixData& pix_data, double& distance)
{
    double dist = 0.0;
    double totalDist = 0.0;

    // We will adjust the minimum distance based on the current zoom
    double epsModified = 0.0;

    int steps = 0;
    vec3 p;
    do
    {
        MULT_SCALAR(p, direction, totalDist);
        ADD_POINT(p, p, from);
        dist = DE(p);

        totalDist += 0.95 * dist;

        epsModified = totalDist;
        epsModified *= eps;
        steps++;
    }
    while (dist > epsModified && totalDist <= render_params.maxDistance &&
        steps < render_params.maxRaySteps);

    if (dist < epsModified)
    {
        // we didnt escape
        pix_data.escaped = false;

        // We hit something, or reached MaxRaySteps
        pix_data.hit = p;

        //figure out the normal of the surface at this point
        vec3 temp;

```

```

    MULT_SCALAR(temp, direction, epsModified);
    SUBTRACT_POINT(temp, p, temp);
    const vec3 normPos = temp;
    normal(normPos, pix_data.normal);
    distance = totalDist;
}
else {
    //we have the background colour
    pix_data.escaped = true;
    distance = 0.0;
}
}
}

```

4.4 renderer.cc

```

/*
This file is part of the Mandelbox program developed for the course
CS/SE Distributed Computer Systems taught by N. Nedialkov in the
Winter of 2015-2016 at McMaster University.

Copyright (C) 2015-2016 T. Gwosdz and N. Nedialkov

This program is free software: you can redistribute it and/or modify
it under the terms of the GNU General Public License as published by
the Free Software Foundation, either version 3 of the License, or
(at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
GNU General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program. If not, see <http://www.gnu.org/licenses/>.
*/
#include <stdio.h>
#include <stdlib.h>
#include "openacc.h"

#include "color.h"
#include "mandelbox.h"
#include "camera.h"
#include "vector3d.h"
#include "3d.h"

#pragma acc routine seq
extern void rayMarch (const RenderParams &render_params, const vec3 &from,
    const vec3 &to, double eps, pixelData &pix_data, double& distance);

#pragma acc routine seq
extern void getColour(vec3 & colour, const pixelData &pixData, const
    RenderParams &render_params, const vec3 &from, const vec3 &direction);

// device pointers declared and allocated in main.cc
extern MandelBoxParams mandelBox_params;
extern vec3* d_to;

```

```

extern vec3* d_colours;
extern double* d_farPoints;
extern pixelData* d_pixData;

// used by main.cc to find new camera position after each frame
extern vec3 newFurthestPoint;

void renderFractal(const CameraParams camera_params, const RenderParams
    renderer_params, unsigned char* image)
{
    // from is const
    vec3 fromTemp;
    VEC(fromTemp, camera_params.camPos[0], camera_params.camPos[1],
        camera_params.camPos[2]);
    const vec3 from = fromTemp;

    const double eps = pow(10.0, renderer_params.detail);

    const int height = renderer_params.height;
    const int width = renderer_params.width;
    const int n = width * height;
    double *distances = (double*)malloc(n * sizeof(double));

    #pragma acc data copyin(camera_params, renderer_params, eps, from)
    #pragma acc data deviceptr(d_to, d_colours, d_farPoints, d_pixData)
    #pragma acc data copyout(image[:n*3], distances[:n])
    {
        #pragma acc kernels loop independent collapse(2)
        for(int j = 0; j < height; j++)
        {
            for(int i = 0; i < width; i++)
            {
                int k = j*width + i;

                UnProject(i, j, camera_params, &(d_farPoints[k*3]));

                SUBTRACT_DARRS(d_to[k], (&(d_farPoints[k*3])), camera_params.camPos);
                NORMALIZE(d_to[k]);

                rayMarch(renderer_params, from, d_to[k], eps, d_pixData[k], distances
[k]);
                getColour(d_colours[k], d_pixData[k], renderer_params, from, d_to[k])
                ;

                //save colour into texture
                image[k*3 + 2] = (unsigned char)(d_colours[k].x * 255);
                image[k*3 + 1] = (unsigned char)(d_colours[k].y * 255);
                image[k*3] = (unsigned char)(d_colours[k].z * 255);
            }
        }
    }

    // find the index of the farthest point
    double maxDistance = 0;
    int maxDistanceIndex = -1;
    for (int i = 0; i < n; i++) {

```

```

    if (distances[i] > maxDistance) {
        maxDistance = distances[i];
        maxDistanceIndex = i;
    }
}

// copy the vector at the furthest point to our new look at
if (maxDistanceIndex >= 0) {
    acc_memcpy_from_device(&newFurthestPoint, &(d_pixData[maxDistanceIndex]
    ].hit), sizeof(vec3));
} else {
    printf("No distance greater than 0 found. Looking at [0,0,0].\n");
    newFurthestPoint.x = 0;
    newFurthestPoint.y = 0;
    newFurthestPoint.z = 0;
}

free(distances);
}

```

4.5 vector3d.h

```

#ifndef vec3_h
#define vec3_h

#ifdef _OPENACC
    #include <acclmath.h>
#else
    #include <math.h>
#endif

typedef struct
{
    double x, y, z;
} vec3;

#define SET_POINT(p,v) { p.x=v[0]; p.y=v[1]; p.z=v[2]; }

#define SUBTRACT_POINT(p,v,u) {\
    p.x=(v.x)-(u.x);\
    p.y=(v.y)-(u.y);\
    p.z=(v.z)-(u.z);\
}

#define SUBTRACT_DARRS(p,d1,d2) {\
    p.x=(d1[0])-(d2[0]);\
    p.y=(d1[1])-(d2[1]);\
    p.z=(d1[2])-(d2[2]);\
}

#define ADD_POINT(p,v,u) {\
    p.x=(v.x)+(u.x);\
    p.y=(v.y)+(u.y);\
    p.z=(v.z)+(u.z);\
}

#define DISTANCE_APART(p,q) ({ sqrt(pow(p.x-q.x, 2) + pow(p.y-q.y, 2) + pow(p

```



```

        .z-q.z, 2)); })

#define NORMALIZE(p) {\
    double fMag = ( p.x*p.x + p.y*p.y + p.z*p.z );\
    if (fMag != 0){\
        double fMult = 1.0/sqrt(fMag);\
        p.x *= fMult;\
        p.y *= fMult;\
        p.z *= fMult;\
    }\
}

#define SUBTRACT_SCALAR(p,v,s) { \
    p.x = v.x - s; \
    p.y = v.y - s; \
    p.z = v.z - s; \
}

#define ADD_SCALAR(p,v,s) {\
    p.x = v.x + s; \
    p.y = v.y + s; \
    p.z = v.z + s; \
}

#define MAGNITUDE(p) ({ sqrt( p.x*p.x + p.y*p.y + p.z*p.z ); })

#define DOT_ASSIGN(d,p) { d=( p.x*p.x + p.y*p.y + p.z*p.z ); }
#define DOT(p,q) { (p.x*q.x + p.y*q.y + p.z*q.z) }
#define MAX(a,b) ( ((a)>(b)) ? (a) : (b))

#define VEC(v,a,b,c) { v.x = a; v.y = b; v.z = c; }

#define MULT_POINTWISE(p,v,u) { p.x = (v.x)*(u.x); p.y = (v.y)*(u.y); p.z = (
    v.z)*(u.z); }

#define MULT_SCALAR(p,v,s) {\
    p.x = v.x * s;\
    p.y = v.y * s;\
    p.z = v.z * s;\
}

#define DIV_SCALAR(p,v,s) {\
    double fInv = 1.0 / s;\
    p.x = v.x * fInv;\
    p.y = v.y * fInv;\
    p.z = v.z * fInv;\
}

inline double clamp(double d, double min, double max)
{
    const double t = d < min ? min : d;
    return t > max ? max : t;
}

inline void clamp(vec3 &v, double min, double max)
{

```

```

    v.x = clamp(v.x,min,max);
    v.y = clamp(v.y,min,max);
    v.z = clamp(v.z,min,max);
}

#endif

```

4.6 3d.cc

```

/*
   This file is part of the Mandelbox program developed for the course
   CS/SE Distributed Computer Systems taught by N. Nedialkov in the
   Winter of 2015-2016 at McMaster University.

   Copyright (C) 2015-2016 T. Gwosdz and N. Nedialkov

   This program is free software: you can redistribute it and/or modify
   it under the terms of the GNU General Public License as published by
   the Free Software Foundation, either version 3 of the License, or
   (at your option) any later version.

   This program is distributed in the hope that it will be useful,
   but WITHOUT ANY WARRANTY; without even the implied warranty of
   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
   GNU General Public License for more details.

   You should have received a copy of the GNU General Public License
   along with this program. If not, see <http://www.gnu.org/licenses/>.
*/

#include <stdlib.h>
#include <string.h>

#define PI 3.14159265358979323846 // M_PI not defined in accelmath
#ifdef _OPENACC
    #include <accelmath.h>
#else
    #include <math.h>
#endif

#include "3d.h"

int UnProject(double winX, double winY, const CameraParams& camP, double *obj
)
{
    //Transformation vectors
    double in[4], out[4];

    //Transformation of normalized coordinates between -1 and 1
    in[0]=(winX-(double)(camP.viewport[0]))/(double)(camP.viewport[2])*2.0-1.0;
    in[1]=(winY-(double)(camP.viewport[1]))/(double)(camP.viewport[3])*2.0-1.0;
    in[2]=2.0-1.0;
    in[3]=1.0;

    //Objects coordinates
    MultiplyMatrixByVector(out, camP.matInvProjModel, in);
}

```

```

    if(out[3]==0.0)
        return 0;

    out[3] = 1.0/out[3];
    obj[0] = out[0]*out[3];
    obj[1] = out[1]*out[3];
    obj[2] = out[2]*out[3];
    return 1;
}

void LoadIdentity(double *matrix){
    matrix[0] = 1.0;
    matrix[1] = 0.0;
    matrix[2] = 0.0;
    matrix[3] = 0.0;

    matrix[4] = 0.0;
    matrix[5] = 1.0;
    matrix[6] = 0.0;
    matrix[7] = 0.0;

    matrix[8] = 0.0;
    matrix[9] = 0.0;
    matrix[10] = 1.0;
    matrix[11] = 0.0;

    matrix[12] = 0.0;
    matrix[13] = 0.0;
    matrix[14] = 0.0;
    matrix[15] = 1.0;
}

void Perspective(double fov, double aspect, double zNear, double zFar, double
    *projMat)
{
    double ymax, xmax;

    ymax = zNear * tan(fov * PI / 360.0);
    xmax = ymax * aspect;
    Frustum(-xmax, xmax, -ymax, ymax, zNear, zFar, projMat);
}

void Frustum(double left, double right, double bottom, double top, double
    znear, double zfar, double *matrix)
{
    double temp, temp2, temp3, temp4;
    temp = 2.0 * znear;
    temp2 = right - left;
    temp3 = top - bottom;
    temp4 = zfar - znear;
    matrix[0] = temp / temp2;
    matrix[1] = 0.0;
    matrix[2] = 0.0;
    matrix[3] = 0.0;
    matrix[4] = 0.0;

```

```

matrix[5] = temp / temp3;
matrix[6] = 0.0;
matrix[7] = 0.0;
matrix[8] = (right + left) / temp2;
matrix[9] = (top + bottom) / temp3;
matrix[10] = (-zfar - znear) / temp4;
matrix[11] = -1.0;
matrix[12] = 0.0;
matrix[13] = 0.0;
matrix[14] = (-temp * zfar) / temp4;
matrix[15] = 0.0;
}
//-----
void LookAt(double *eye, double *target, double *upV, double *modelMatrix)
{
    double forward[3], side[3], up[3];
    double matrix2[16], resultMatrix[16];
    //-----
    forward[0] = target[0] - eye[0];
    forward[1] = target[1] - eye[1];
    forward[2] = target[2] - eye[2];
    NormalizeVector(forward);
    //-----
    //Side = forward x up
    ComputeNormalOfPlane(side, forward, upV);
    NormalizeVector(side);
    //-----
    //Recompute up as: up = side x forward
    ComputeNormalOfPlane(up, side, forward);
    //-----
    matrix2[0] = side[0];
    matrix2[4] = side[1];
    matrix2[8] = side[2];
    matrix2[12] = 0.0;
    //-----
    matrix2[1] = up[0];
    matrix2[5] = up[1];
    matrix2[9] = up[2];
    matrix2[13] = 0.0;
    //-----
    matrix2[2] = -forward[0];
    matrix2[6] = -forward[1];
    matrix2[10] = -forward[2];
    matrix2[14] = 0.0;
    //-----
    matrix2[3] = matrix2[7] = matrix2[11] = 0.0;
    matrix2[15] = 1.0;
    //-----
    MultiplyMatrices(resultMatrix, modelMatrix, matrix2);
    Translate(resultMatrix, -eye[0], -eye[1], -eye[2]);
    //-----
    memcpy(modelMatrix, resultMatrix, 16*sizeof(double));
}

void NormalizeVector(double *v)

```

```

{
    double m = 1.0/sqrt(v[0]*v[0]+v[1]*v[1]+v[2]*v[2]);
    v[0] *= m;
    v[1] *= m;
    v[2] *= m;
}

void ComputeNormalOfPlane(double *normal, double *v1, double *v2)
{
    normal[0] = v1[1] * v2[2] - v1[2] * v2[1];
    normal[1] = v1[2] * v2[0] - v1[0] * v2[2];
    normal[2] = v1[0] * v2[1] - v1[1] * v2[0];
}

void MultiplyMatrices(double *result, const double *matrix1, const double *
    matrix2)
{
    result[0]=matrix1[0]*matrix2[0]+
        matrix1[4]*matrix2[1]+
        matrix1[8]*matrix2[2]+
        matrix1[12]*matrix2[3];
    result[4]=matrix1[0]*matrix2[4]+
        matrix1[4]*matrix2[5]+
        matrix1[8]*matrix2[6]+
        matrix1[12]*matrix2[7];
    result[8]=matrix1[0]*matrix2[8]+
        matrix1[4]*matrix2[9]+
        matrix1[8]*matrix2[10]+
        matrix1[12]*matrix2[11];
    result[12]=matrix1[0]*matrix2[12]+
        matrix1[4]*matrix2[13]+
        matrix1[8]*matrix2[14]+
        matrix1[12]*matrix2[15];
    result[1]=matrix1[1]*matrix2[0]+
        matrix1[5]*matrix2[1]+
        matrix1[9]*matrix2[2]+
        matrix1[13]*matrix2[3];
    result[5]=matrix1[1]*matrix2[4]+
        matrix1[5]*matrix2[5]+
        matrix1[9]*matrix2[6]+
        matrix1[13]*matrix2[7];
    result[9]=matrix1[1]*matrix2[8]+
        matrix1[5]*matrix2[9]+
        matrix1[9]*matrix2[10]+
        matrix1[13]*matrix2[11];
    result[13]=matrix1[1]*matrix2[12]+
        matrix1[5]*matrix2[13]+
        matrix1[9]*matrix2[14]+
        matrix1[13]*matrix2[15];
    result[2]=matrix1[2]*matrix2[0]+
        matrix1[6]*matrix2[1]+
        matrix1[10]*matrix2[2]+
        matrix1[14]*matrix2[3];
    result[6]=matrix1[2]*matrix2[4]+
        matrix1[6]*matrix2[5]+
        matrix1[10]*matrix2[6]+

```

```

    matrix1[14]*matrix2[7];
result[10]=matrix1[2]*matrix2[8]+
    matrix1[6]*matrix2[9]+
    matrix1[10]*matrix2[10]+
    matrix1[14]*matrix2[11];
result[14]=matrix1[2]*matrix2[12]+
    matrix1[6]*matrix2[13]+
    matrix1[10]*matrix2[14]+
    matrix1[14]*matrix2[15];
result[3]=matrix1[3]*matrix2[0]+
    matrix1[7]*matrix2[1]+
    matrix1[11]*matrix2[2]+
    matrix1[15]*matrix2[3];
result[7]=matrix1[3]*matrix2[4]+
    matrix1[7]*matrix2[5]+
    matrix1[11]*matrix2[6]+
    matrix1[15]*matrix2[7];
result[11]=matrix1[3]*matrix2[8]+
    matrix1[7]*matrix2[9]+
    matrix1[11]*matrix2[10]+
    matrix1[15]*matrix2[11];
result[15]=matrix1[3]*matrix2[12]+
    matrix1[7]*matrix2[13]+
    matrix1[11]*matrix2[14]+
    matrix1[15]*matrix2[15];
}

inline void MultiplyMatrixByVector(double *resultvector, const double *matrix
, double *pvector)
{
    resultvector[0]=matrix[0]*pvector[0]+matrix[4]*pvector[1]+matrix[8]*pvector
    [2]+matrix[12]*pvector[3];
    resultvector[1]=matrix[1]*pvector[0]+matrix[5]*pvector[1]+matrix[9]*pvector
    [2]+matrix[13]*pvector[3];
    resultvector[2]=matrix[2]*pvector[0]+matrix[6]*pvector[1]+matrix[10]*
    pvector[2]+matrix[14]*pvector[3];
    resultvector[3]=matrix[3]*pvector[0]+matrix[7]*pvector[1]+matrix[11]*
    pvector[2]+matrix[15]*pvector[3];
}

#define SWAP_ROWS(a, b) { double *_tmp = a; (a)=(b); (b)=_tmp; }
#define MAT(m,r,c) (m)[(c)*4+(r)]

int InvertMatrix(double *m, double *out){
    double wtmp[4][8];
    double m0, m1, m2, m3, s;
    double *r0, *r1, *r2, *r3;
    r0 = wtmp[0], r1 = wtmp[1], r2 = wtmp[2], r3 = wtmp[3];
    r0[0] = MAT(m, 0, 0), r0[1] = MAT(m, 0, 1),
    r0[2] = MAT(m, 0, 2), r0[3] = MAT(m, 0, 3),
    r0[4] = 1.0, r0[5] = r0[6] = r0[7] = 0.0,
    r1[0] = MAT(m, 1, 0), r1[1] = MAT(m, 1, 1),
    r1[2] = MAT(m, 1, 2), r1[3] = MAT(m, 1, 3),
    r1[4] = 1.0, r1[5] = r1[6] = r1[7] = 0.0,
    r2[0] = MAT(m, 2, 0), r2[1] = MAT(m, 2, 1),
    r2[2] = MAT(m, 2, 2), r2[3] = MAT(m, 2, 3),

```

```

    r2[6] = 1.0, r2[4] = r2[5] = r2[7] = 0.0,
    r3[0] = MAT(m, 3, 0), r3[1] = MAT(m, 3, 1),
    r3[2] = MAT(m, 3, 2), r3[3] = MAT(m, 3, 3),
    r3[7] = 1.0, r3[4] = r3[5] = r3[6] = 0.0;
/* choose pivot - or die */
if (fabs(r3[0]) > fabs(r2[0]))
    SWAP_ROWS(r3, r2);
if (fabs(r2[0]) > fabs(r1[0]))
    SWAP_ROWS(r2, r1);
if (fabs(r1[0]) > fabs(r0[0]))
    SWAP_ROWS(r1, r0);
if (0.0 == r0[0])
    return 0;
/* eliminate first variable */
m1 = r1[0] / r0[0];
m2 = r2[0] / r0[0];
m3 = r3[0] / r0[0];
s = r0[1];
r1[1] -= m1 * s;
r2[1] -= m2 * s;
r3[1] -= m3 * s;
s = r0[2];
r1[2] -= m1 * s;
r2[2] -= m2 * s;
r3[2] -= m3 * s;
s = r0[3];
r1[3] -= m1 * s;
r2[3] -= m2 * s;
r3[3] -= m3 * s;
s = r0[4];
if (s != 0.0) {
    r1[4] -= m1 * s;
    r2[4] -= m2 * s;
    r3[4] -= m3 * s;
}
s = r0[5];
if (s != 0.0) {
    r1[5] -= m1 * s;
    r2[5] -= m2 * s;
    r3[5] -= m3 * s;
}
s = r0[6];
if (s != 0.0) {
    r1[6] -= m1 * s;
    r2[6] -= m2 * s;
    r3[6] -= m3 * s;
}
s = r0[7];
if (s != 0.0) {
    r1[7] -= m1 * s;
    r2[7] -= m2 * s;
    r3[7] -= m3 * s;
}
/* choose pivot - or die */
if (fabs(r3[1]) > fabs(r2[1]))
    SWAP_ROWS(r3, r2);

```

```

if (fabs(r2[1]) > fabs(r1[1]))
    SWAP_ROWS(r2, r1);
if (0.0 == r1[1])
    return 0;
/* eliminate second variable */
m2 = r2[1] / r1[1];
m3 = r3[1] / r1[1];
r2[2] -= m2 * r1[2];
r3[2] -= m3 * r1[2];
r2[3] -= m2 * r1[3];
r3[3] -= m3 * r1[3];
s = r1[4];
if (0.0 != s) {
    r2[4] -= m2 * s;
    r3[4] -= m3 * s;
}
s = r1[5];
if (0.0 != s) {
    r2[5] -= m2 * s;
    r3[5] -= m3 * s;
}
s = r1[6];
if (0.0 != s) {
    r2[6] -= m2 * s;
    r3[6] -= m3 * s;
}
s = r1[7];
if (0.0 != s) {
    r2[7] -= m2 * s;
    r3[7] -= m3 * s;
}
/* choose pivot - or die */
if (fabs(r3[2]) > fabs(r2[2]))
    SWAP_ROWS(r3, r2);
if (0.0 == r2[2])
    return 0;
/* eliminate third variable */
m3 = r3[2] / r2[2];
r3[3] -= m3 * r2[3], r3[4] -= m3 * r2[4],
    r3[5] -= m3 * r2[5], r3[6] -= m3 * r2[6], r3[7] -= m3 * r2[7];
/* last check */
if (0.0 == r3[3])
    return 0;
s = 1.0 / r3[3];      /* now back substitute row 3 */
r3[4] *= s;
r3[5] *= s;
r3[6] *= s;
r3[7] *= s;
m2 = r2[3];          /* now back substitute row 2 */
s = 1.0 / r2[2];
r2[4] = s * (r2[4] - r3[4] * m2), r2[5] = s * (r2[5] - r3[5] * m2),
    r2[6] = s * (r2[6] - r3[6] * m2), r2[7] = s * (r2[7] - r3[7] * m2);
m1 = r1[3];
r1[4] -= r3[4] * m1, r1[5] -= r3[5] * m1,
    r1[6] -= r3[6] * m1, r1[7] -= r3[7] * m1;
m0 = r0[3];

```



```

    r0[4] -= r3[4] * m0, r0[5] -= r3[5] * m0,
    r0[6] -= r3[6] * m0, r0[7] -= r3[7] * m0;
    m1 = r1[2]; /* now back substitute row 1 */
    s = 1.0 / r1[1];
    r1[4] = s * (r1[4] - r2[4] * m1), r1[5] = s * (r1[5] - r2[5] * m1),
    r1[6] = s * (r1[6] - r2[6] * m1), r1[7] = s * (r1[7] - r2[7] * m1);
    m0 = r0[2];
    r0[4] -= r2[4] * m0, r0[5] -= r2[5] * m0,
    r0[6] -= r2[6] * m0, r0[7] -= r2[7] * m0;
    m0 = r0[1]; /* now back substitute row 0 */
    s = 1.0 / r0[0];
    r0[4] = s * (r0[4] - r1[4] * m0), r0[5] = s * (r0[5] - r1[5] * m0),
    r0[6] = s * (r0[6] - r1[6] * m0), r0[7] = s * (r0[7] - r1[7] * m0);
    MAT(out, 0, 0) = r0[4];
    MAT(out, 0, 1) = r0[5], MAT(out, 0, 2) = r0[6];
    MAT(out, 0, 3) = r0[7], MAT(out, 1, 0) = r1[4];
    MAT(out, 1, 1) = r1[5], MAT(out, 1, 2) = r1[6];
    MAT(out, 1, 3) = r1[7], MAT(out, 2, 0) = r2[4];
    MAT(out, 2, 1) = r2[5], MAT(out, 2, 2) = r2[6];
    MAT(out, 2, 3) = r2[7], MAT(out, 3, 0) = r3[4];
    MAT(out, 3, 1) = r3[5], MAT(out, 3, 2) = r3[6];
    MAT(out, 3, 3) = r3[7];
    return 1;
}

void Translate(double *result, double x, double y, double z){
    double matrix[16], resultMatrix[16];

    LoadIdentity(matrix);
    matrix[12] = x;
    matrix[13] = y;
    matrix[14] = z;

    MultiplyMatrices(resultMatrix, result, matrix);
    memcpy(result, resultMatrix, 16*sizeof(double));
}

```

4.7 3d.h

```

/*
This file is part of the Mandelbox program developed for the course
CS/SE Distributed Computer Systems taught by N. Nediaklov in the
Winter of 2015-2016 at McMaster University.

Copyright (C) 2015-2016 T. Gwosdz and N. Nediaklov

This program is free software: you can redistribute it and/or modify
it under the terms of the GNU General Public License as published by
the Free Software Foundation, either version 3 of the License, or
(at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
GNU General Public License for more details.

```

```

    You should have received a copy of the GNU General Public License
    along with this program.  If not, see <http://www.gnu.org/licenses/>.
*/

#ifdef _3d_H
#define _3d_H

#define NEAR 1
#define FAR 100

#include "camera.h"
#include "renderer.h"

void LoadIdentity (double *matrix);
void Perspective (double fov, double aspect, double zNear, double zFar,
    double *projMatrix);
void Frustum (double left, double right, double bottom, double top,
    double znear, double zfar, double *matrix);
void LookAt (double *eye, double *target, double *up, double *
    modelMatrix);
double LengthVector (double *vector);
void NormalizeVector(double *vector);
void ComputeNormalOfPlane(double *normal, double *v1, double *v2);
void MultiplyMatrices(double *result, const double *matrix1, const double *
    matrix2);
void MultiplyMatrixByVector(double *resultvector, const double *matrix,
    double *pvector);
int InvertMatrix(double *m, double *out);
void Translate(double *result, double x, double y, double z);
int UnProject(double winX, double winY, double winZ, const double *model,
    const double *proj, const int *view, double *obj);

#pragma acc routine seq
int UnProject(double winX, double winY, const CameraParams & camP, double
    *obj);

#endif

```

5 Running the Program

To generate the video, follow the following steps.

1. Run `make mandelbox` from the project directory
2. Execute the program with `./mandelbox params_mandelbox.dat`
3. Convert the images to a video by running `./video.sh`

If you wish to change the number of frames, simply open `main.cc` and replace the value for the `NUM_FRAMES` pragma with the desired number of frames.

6 Bonus Features

6.1 Automatic Navigation

The automatic navigation functionality for the project was implemented. That is, the program will determine a path through the mandelbox that does not “hit” any walls, nor leave the box itself. For more details on the algorithm used, refer to section 2.